


GPU-Based Ray-Casting

Scientific Visualization

Pere-Pau Vázquez – Dep. Computer Science – UPC



Outline

- Ray-casting of rectilinear (structured) grids
- Optimizations




Why Ray-Casting on GPUs?

- Most GPU rendering is object-order (rasterization)
- Image-order is more “CPU-like”
 - Recent fragment shader advances
 - Simpler to implement
 - Very flexible (e.g., adaptive sampling)
 - Correct perspective projection
 - Can be implemented in single pass!
 - Native 32-bit compositing




Where Is Correct Perspective Needed?

- Entering the volume
- Wide field of view
- Fly-throughs
- Virtual endoscopy
- Integration into perspective scenes, e.g., games




Recent GPU Ray-Casting Approaches

- Rectilinear grids
 - [Krüger and Westermann, 2003]
 - [Röttger et al., 2003]
 - [Green, 2004] (NVIDIA SDK Example)
 - [Stegmaier et al., 2005]
 - [Scharsach et al., 2006]
- Unstructured (tetrahedral) grids
 - [Weiler et al., 2002, 2003, 2004]
 - [Bernardon, 2004]




Single-Pass Ray-Casting

- Enabled by conditional loops in fragment shaders (since Shader Model 3; e.g., Geforce 6800, ATI X1800)
- Substitute multiple passes and early-z testing by single loop and early loop exit
- No compositing buffer: full 32-bit precision!
- NVIDIA example: compute ray intersections with bounding box, march along rays and composite




Basic Ray Setup / Termination

- Two main approaches:
 - Procedural ray/box intersection [Röttger et al., 2003], [Green, 2004]
 - Rasterize bounding box [Krüger and Westermann, 2003]
- Some possibilities
 - Ray start position and exit check
 - Ray start position and exit position
 - Ray start position and direction vector




Parallel volume rendering with a single GPU

- Many volume rendering algorithms cannot fully exploit parallel pipelines
 - Cell projection: requires visibility sorting of cells
 - Textured slices: rasterize all fragments of all slices



Parallel volume rendering with a single GPU

- Ray casting can exploit parallel pixel pipelines:
 - Like ray tracing, ray casting is embarrassingly parallel




Parallel volume rendering with a single GPU

- More benefits of ray casting on a GPU:
 - Applicable to uniform and tetrahedral meshes
 - Important optimizations can be implemented
 - We will see later...


Parallel volume rendering with a single GPU

- Basic idea
 - Multi-pass approach: Render screen-filling rectangles to call a program for each pixel




Parallel volume rendering with a single GPU

- Initialize pixels with first intersection of the rays with the volume data




Parallel volume rendering with a single GPU

- In each pass, propagate view rays and store accumulated color and sampling position




Parallel volume rendering with a single GPU

- In each pass, propagate view rays and store accumulated color and sampling position



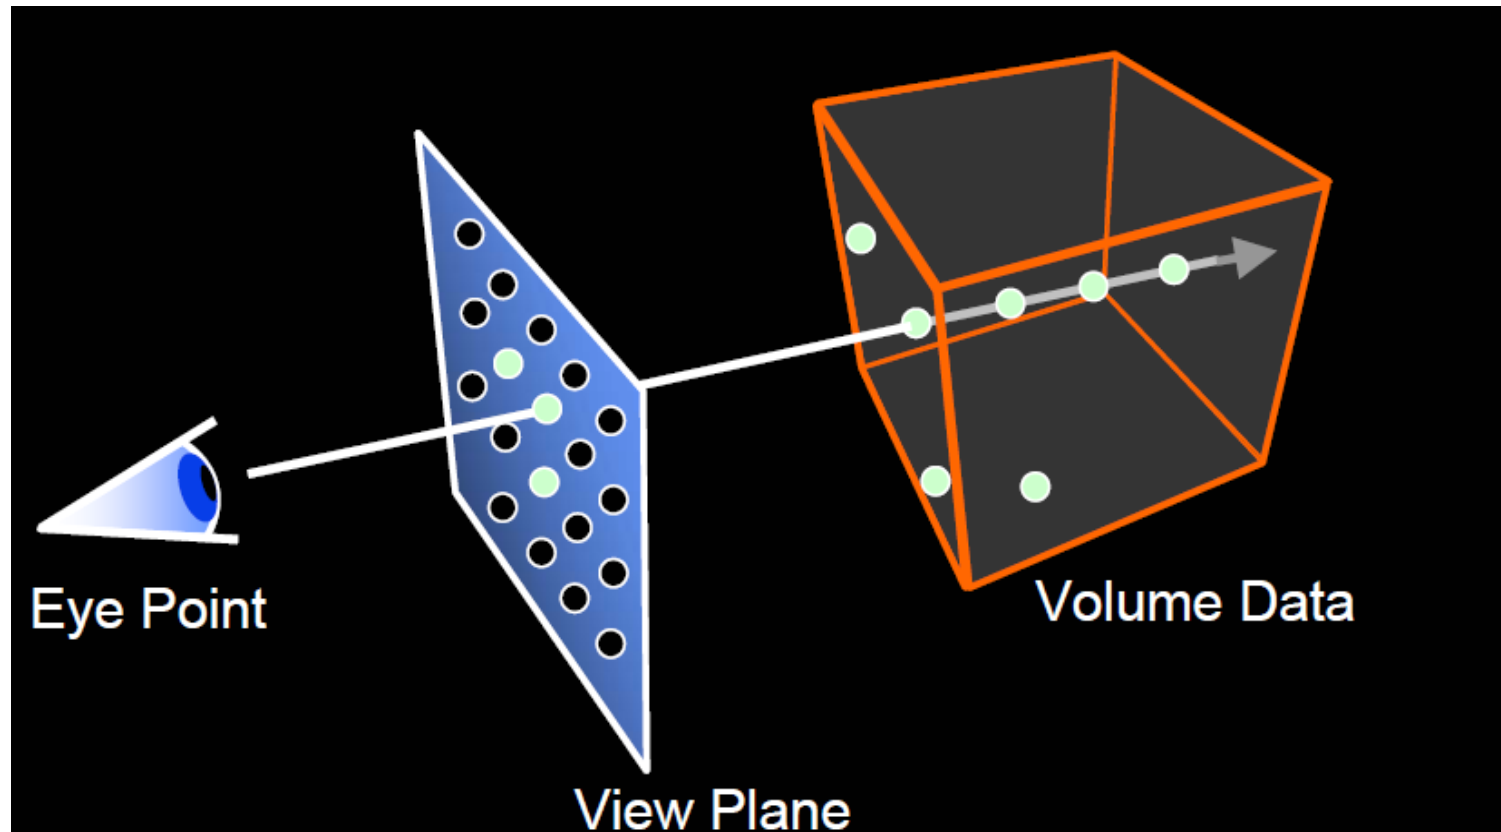
Parallel volume rendering with a single GPU

- In each pass, propagate view rays and store accumulated color and sampling position



Parallel volume rendering with a single GPU

- Stop when all rays have left the volume



Parallel volume rendering with a single GPU.

Optimizations

- Avoid work for rays that don't intersect the volume
- Avoid work for rays that have accumulated full opacity
- Quickly test whether all rays have left the volume
- Quickly go through empty regions
- Adapt sampling distance to data


Parallel volume rendering with a single GPU.

Optimizations

- **Avoid work for rays that don't intersect the volume**
- Avoid work for rays that have accumulated full opacity
- Quickly test whether all rays have left the volume
- Quickly go through empty regions
- Adapt sampling distance to data


Avoid rays outside the volume

- Use rasterization to start rays
 - Render bbox & issue fragment shader on valid fragments



Procedural Ray Setup/Termination

- Everything handled in the fragment shader
- Procedural ray / bounding box intersection
- Ray is given by camera position and volume entry position
- Exit criterion needed
- Pro: simple and self-contained
- Con: full load on the fragment shader





Fragment Shader

- Rasterize front faces of volume bounding box
- Texcoords are volume position in [0,1]
- Subtract camera position
- Repeatedly check for exit of bounding box

```
// Cg fragment shader code for single-pass ray casting
float4 main(VS_OUTPUT IN, float4 TexCoord0 : TEXCOORD0,
            uniform sampler3D SamplerDataVolume,
            uniform sampler1D SamplerTransferFunction,
            uniform float3 camera,
            uniform float stepsize,
            uniform float3 volExtentMin,
            uniform float3 volExtentMax
            ) : COLOR
{
    float4 value;
    float scalar;
    // Initialize accumulated color and opacity
    float4 dst = float4(0,0,0,0);
    // Determine volume entry position
    float3 position = TexCoord0.xyz;
    // Compute ray direction
    float3 direction = TexCoord0.xyz - camera;
    direction = normalize(direction);
    // Loop for ray traversal
    for (int i = 0; i < 200; i++) // Some large number
    {
        // Data access to scalar value in 3D volume texture
        value = tex3D(SamplerDataVolume, position);
        scalar = value.a;
        // Apply transfer function
        float4 src = tex1D(SamplerTransferFunction, scalar);
        // Front-to-back compositing
        dst = (1.0-dst.a) * src + dst;
        // Advance ray position along ray direction
        position = position + direction * stepsize;
        // Ray termination: Test if outside volume ...
        float3 temp1 = sign(position - volExtentMin);
        float3 temp2 = sign(volExtentMax - position);
        float inside = dot(temp1, temp2);
        // ... and exit loop
        if (inside < 3.0)
            break;
    }
    return dst;
}
```

"Image-Based" Ray Setup/Termination

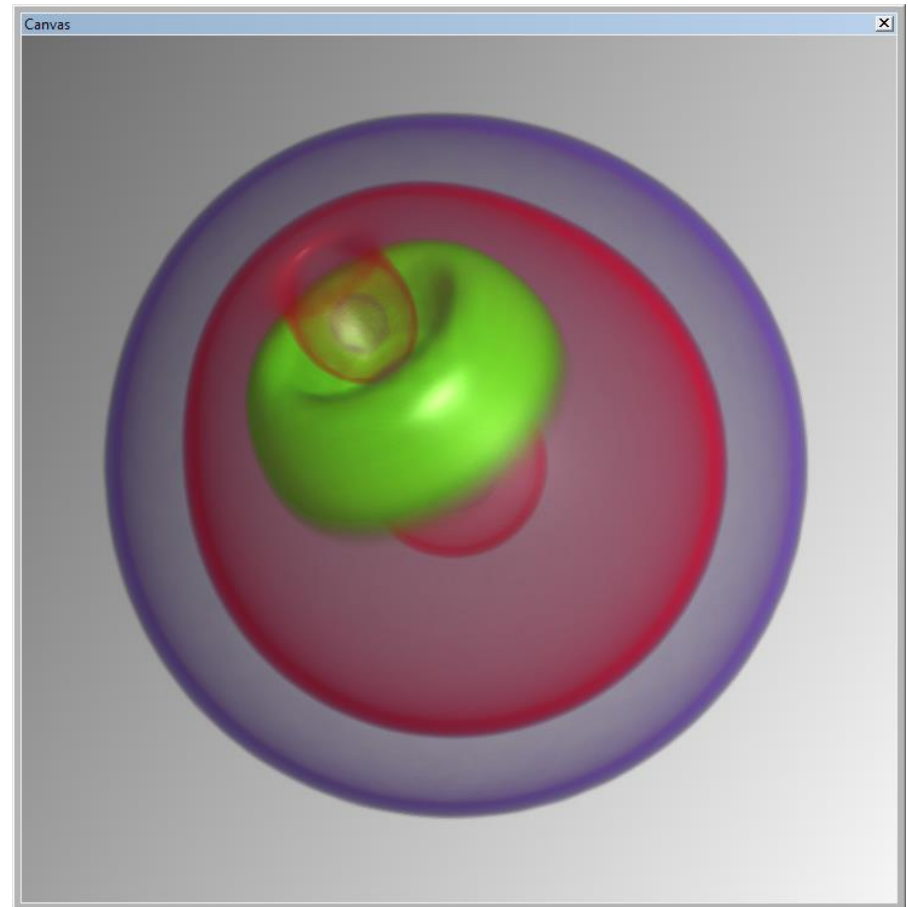

- Rasterize bounding box front faces and back faces [Krüger and Westermann, 2003]
- Ray start position: front faces
- Direction vector: back-front faces



- Independent of projection (orthogonal/perspective)



Is all this perfect?

- Need interpolation hardware (or software)





Is all this perfect?

- No, it is not. Sampling rate must be high enough.




Is all this perfect?

- Need to compute gradients to have a nice lighting solution





Is all this perfect?

- No, it is not.



Is all this perfect?

- No, it is not.




Parallel volume rendering with a single GPU. Optimizations

- *Avoid work for rays that don't intersect the volume*
- **Avoid work for rays that have accumulated full opacity**
- Quickly test whether all rays have left the volume
- Quickly go through empty regions
- Adapt sampling distance to data

Standard Ray-Casting Optimizations (1)

Early ray termination

- Isosurfaces: stop when surface hit
- Direct volume rendering:
stop when opacity \geq threshold
- Several possibilities
 - Older GPUs: multi-pass rendering with early-z test
 - Shader model 3: break out of ray-casting loop
 - Current GPUs: early loop exit not optimal but good



Parallel volume rendering with a single GPU.


Optimizations

- *Avoid work for rays that don't intersect the volume*
- *Avoid work for rays that have accumulated full opacity*
- Quickly test whether all rays have left the volume
- **Quickly go through empty regions**
- Adapt sampling distance to data


Standard Ray-Casting Optimizations (2)

Empty space skipping


- Skip transparent samples
 - Depends on transfer function
 - Start casting close to first hit
-
- Several possibilities
 - Per-sample check of opacity (expensive)
 - Traverse hierarchy (e.g., octree) or regular grid
 - These are image-order: what about object-order?



Object-Order Empty Space Skipping (1)




rasterize bounding box



rasterize "tight" bounding geometry


Object-Order Empty Space Skipping (2)

- Store min-max values of volume bricks
- Cull bricks against isovalue or transfer function
- Rasterize front and back faces of active bricks




Object-Order Empty Space Skipping (3)

- Rasterize front and back faces of active min-max bricks
- Start rays on brick front faces
- Terminate when
 - Full opacity reached, or
 - Back face reached



Object-Order Empty Space Skipping (3)

- Rasterize front and back faces of active min-max bricks
- Start rays on brick front faces
- Terminate when
 - Full opacity reached, or
 - Back face reached
- Not all empty space is skipped



Isosurface Ray-Casting

- Isosurfaces/Level Sets


- scanned data
- distance fields
- CSG operations
- level sets: surface editing, simulation, segmentation, ...



Intersection Refinement (1)

- Fixed number of bisection or binary search steps
- Virtually no impact on performance

- Refine already detected intersection
- Handle problems with small features / at silhouettes with adaptive sampling



Intersection Refinement (2)

without refinement

with refinement



sampling rate $1/5$ voxel (no adaptive sampling)

Intersection Refinement (3)




Sampling distance 1.0, 24 fps



Sampling distance 5.0, 66 fps


Deferred Isosurface Shading

- Shading is expensive
 - Gradient computation; conditional execution not free
- Ray-casting step computes only intersection image




Enhancements (1)

- Build on image-based ray setup
- Allow viewpoint inside the volume



- Intersect polygonal geometry



Enhancements (2)

1. Starting position computation

⇒ Ray start position image

2. Ray length computation

⇒ Ray length image

3. Render polygonal geometry


⇒ Modified ray length image

4. Raycasting

⇒ Compositing buffer



5. Blending

⇒ Final image



Moving Into The Volume (1)

- Near clipping plane clips into front faces



- Fill in holes with near clipping plane
- Can use depth buffer [Scharsach et al., 2006]

Moving Into The Volume (2)

1. Rasterize near clipping plane


- Disable depth buffer, enable color buffer
- Rasterize entire near clipping plane

2. Rasterize nearest back faces

- Enable depth buffer, disable color buffer
- Rasterize *nearest back faces* of active bricks

3. Rasterize nearest front faces

- Enable depth buffer, enable color buffer
- Rasterize *nearest front faces* of active bricks



How to speed-up?


- Algorithmic improvements
 - Many opportunities
- More hardware
 - Clusters of computers
 - Client-server approach
- Better hardware
 - Specialized hardware?

Raycasting pipeline

Generate rays for pixels

```
While (in volume){  
    read data from volume  
    map data to color  
    accumulate color  
}
```

Write color to framebuffer



Raycasting pipeline

Generate rays for pixels

Generate less rays

```
While (in volume){  
    read data from volume  
    map data to color  
    accumulate color  
}
```

Write color to framebuffer

Raycasting pipeline

Generate rays for pixels

While (in volume){

 use fewer samples

 read data from volume

 map data to color

 accumulate color

}

Write color to framebuffer

Raycasting pipeline

Generate rays for pixels

While (in volume){

 read data from volume

 read data faster

 map data to color

 accumulate color

}

Write color to framebuffer

Raycasting pipeline

Generate rays for pixels

While (in volume){

 read data from volume

 map data to color

 map less often

 accumulate color

}

Write color to framebuffer

Raycasting pipeline

Generate rays for pixels

While (in volume){

 read data from volume

 map data to color

 accumulate color

 accumulate only if needed

}

Write color to framebuffer

Generate less rays



- Do we need one ray per pixel?
 - Often not!
- Shoot rays in interesting regions
 - How to identify them?
 - Adaptive image sampling

Generate less rays

- Tile-based sampling by Ljung
 - Split the image to fixed-size tiles
 - Balance overhead vs. gains
 - Shoot probe rays to evaluate tile importance
 - Pre-classified volume blocks importance
 - Assign tile sampling rate
 - Higher importance -> more samples in tile
 - Not restricted to pixel-aligned samples
 - Arbitrary sampling within the tile

Use fewer samples

- Only about 5-10% of voxels contribute to image
 - Why to traverse all?
 - Depends on transfer function!



Images by [Marmitt 06]

Skip empty regions

- Block-based skipping
 - Encode empty space in additional structure
 - Trade speed for space!
 - Octrees & kD-trees
 - Implicit tree with min/max values in each node
 - Summed-area table for transfer function
 - Both semi-transparent and isosurfaces
 - Size may be a problem
 - Fat blocks
 - Fixed-size block with min/max values
 - Limited effectivity

Read data faster

- Volume swizzling
 - Heavy memory bandwidth requirements
 - Raycasting generates scattered memory access
 - Cache unfriendly
 - Maps neighboring pixels to be close in memory
 - Increases cache efficiency
 - Noticeable speed-up (2-3x)

Map less often

- For expensive shading
 - Complex local lighting models
 - Shadows, Fresnel approximation, ambient occlusion
- Shade only if required
- Skip highly transparent samples
 - Contribution is insignificant
- Skip low gradient magnitude samples
 - Improves more quality than speed

Accumulate only if needed

- Early ray termination
 - Most samples are hidden behind opaque parts
 - Stop the computation if the ray is saturated
 - Empirical threshold = 0.95
 - Trivial implementation

Considerations on GPU-based ray-casting

- Most flexible algorithm
 - Can incorporate speed-up techniques easily
- Balances the requirements on the GPU
 - Can use texture interpolation hardware
 - Can use rasterizers, but **does not depend** on them
- Shader-based vs. CUDA-based implementations
 - On par

Considerations on GPU-based ray-casting

- Transfer function implemented as a texture
 - May require more complex systems for sophisticated transfer functions
- Empty space leaping
 - May require extremely high amounts of storage to make it efficient
- May use multiple-pass algorithms
 - E.g. for deferred shading

Acknowledgments

- Built from slides by Henning Scharsach, Christian Sigg, Daniel Weiskopf, Martin Kraus, Lukas Marsalek, Timo Ropinski, Christof Rezk-Salama, Klaus Engel, Markus Hadwiger... Thanks to all of them!!!
- Check www.real-time-volume-graphics.org for more materials